

# Reinforcement learning for symbolic mathematics

Kevin, O’Keeffe<sup>\*1</sup>

## Abstract

We explore whether reinforcement learning (RL) can be useful for symbolic mathematics. Previous work showed contrastive learning can solve linear equations in one variable. We show that model-free PPO (Schulman et al., 2017) combined with a graph-based policy (TreeMLP), a dynamic per-state action space, and lightweight inductive biases (curriculum, success replay, constant relabeling) can solve nonlinear equations involving radicals, exponentials, and trigonometric functions—reaching or exceeding state-of-the-art accuracy on the CommonCore benchmark. We further consider *open* equations whose solution requires a generative change-of-variables, and train a separate substitution policy  $\pi_{\text{cov}}$  to handle them.




Figure 1. Expression trees for  $ax + b$  and  $ax^2 + bx + c$ .

## 1. Introduction

Reinforcement learning (RL) has been applied to diverse fields (Ng et al., 2006; Degraeve et al., 2022; Vinyals et al., 2019) but has not yet been widely adopted in symbolic mathematics, where agents perform tasks like solving algebraic equations or evaluating integrals analytically. Traditional approaches to symbolic mathematics rely on hand-engineered systems like Mathematica or Maple. An RL-based approach, where agents learn transformations autonomously, could reduce manual curation and find novel solution techniques.

Applying RL to symbolic math is difficult because the state space is combinatorially large, as each equation can branch into numerous sub-expressions. The action space is also large and dynamic: at every step, the agent can apply various algebraic manipulations to different sub-expressions, so the number and type of actions change with the equation’s form.

We show that RL agents combined with a graph-based policy (TreeMLP), a dynamic per-state action space, and lightweight inductive biases (curriculum learning, success

replay, constant relabeling) can overcome these challenges. The agents learn to solve a wide range of algebraic equations including those involving radicals, exponentials, and trigonometric functions. Although perhaps simple from a human perspective, solving such equations with RL is non-trivial and—to our knowledge—has not been done before. Symbolic equation solving is also a key building block for more advanced tasks like solving differential equations, the natural next step.

## 2. Problem Formulation

Our Markov Decision Process (MDP) formulation is

**States:** equations like  $ax + b = 0$  or  $cx + d = -x/b$ . We vectorize these using preorder traversal over the equation’s expression tree (Figure 1) and pad to a maximum length  $L = 50$  (long enough for the equation sizes we consider). Operations and symbols are mapped to integers, e.g.  $\{\text{add} : 1, \text{sub} : 2, \text{mul} : 3, \dots, x : 5, a : 6, b : 7, \dots\}$ . We encode both the lhs and rhs of the equation in this way and concatenate. For example,  $x + a \rightarrow [\text{add}, x, a] \rightarrow [1, 5, 6, \text{PAD}, \dots]$ .

**Actions:** represented as  $(\text{operation}, \text{term})$  pairs, such as

<sup>\*</sup>Equal contribution <sup>1</sup>Starling Research Institute, Seattle, USA. Correspondence to: Kevin, O’Keeffe <kevin.p.okeeffe@gmail.com>.

$(\text{sub}, b)$  or  $(\text{div}, a)$ . Formally,

$$O = \{\text{add}, \text{sub}, \text{mul}, \text{div}\} \cup \{\text{square}, \text{sqrt}, \text{exp}, \text{log}, \text{sin}, \text{cos}, \text{asin}, \text{acos}\}, \quad (1)$$

$$T = \text{SubExpr}(\text{lhs}) \cup \text{SubExpr}(\text{rhs}), \quad (2)$$

$$A = (O \times T) \cup \{(\text{expand}, \text{None}), (\text{collect}, x), (\text{mul}, -1)\}. \quad (3)$$

Notice the first set of operations take two arguments, the second one argument. For the terms, we choose the list of sub-expressions in the lhs and rhs expression trees. Example of sub-expressions are

$$ax + b = 0 \Rightarrow \{a, x, ax, b\} \quad (4)$$

$$\frac{ax + b}{cx + d} + e = 0 \Rightarrow \{a, x, ax, b, c, d, cx, cx + d, ax + b\} \quad (5)$$

This term set is expressive enough to solve rational equations and all other equation types we consider in this paper. It is also dynamic: the list of sub-expressions is derived from the equation/state and thus has variable length. Looping back to the operations, we also include  $(\text{mul}, -1)$  and

$$\begin{aligned} \text{expand: } & cx + d + e(ax + b) \rightarrow aex + be + cx + d \\ \text{collect } x: & ax + bx \rightarrow (a + b)x \end{aligned}$$

These allow the agent to perform key algebraic steps (Appendix). Finally, we index the action set  $A$  serially, cap it at size  $|A| = 50^1$  and mask illegal actions (e.g., division by zero, see Appendix).

**Rewards.** Define the *complexity*  $C$  of an equation as the total number of nodes and edges in the expression tree <sup>2</sup>. The complexities of the equations in Figure 1 are  $C(ax + b) = 5 + 4 = 9$  and  $C(ax^2 + bx + c) = 10 + 9 = 19$ . The reward is then

$$R(\text{action}) = C(\text{equation}) - C(\text{equation after action}) \quad (6)$$

The intuition here is to encourage the agent to take actions which simplify the equation.

**State Transition Function.** We wrote our code in Python and used SymPy to apply operations to the equations. At each step, we keep track of a lhs and rhs of an equation e.g.  $\text{lhs} = (ax + b)$  and  $\text{rhs} = 0$ . We apply actions to both the lhs and rhs. For example,  $(\text{sub}, b)$  results in  $(\text{lhs}, \text{rhs}) = (ax, -b)$ , and then  $(\text{div}, a)$  results

in  $(\text{lhs}, \text{rhs}) = (x, -b/a)$ . The terminal condition for the environment is when  $\text{lhs} = x$  and the substitution of the  $\text{rhs}$  into the original equation simplifies to 0.

**Limitations.** Importantly, this MDP formulation only works on equations which are ‘closed,’ in the sense that solving them requires manipulating the terms already present in the equation/sub-expression list. By contrast, solving ‘open’ equations requires adding new, out-of-equation terms or clever substitutions. A classic example is the quadratic equation  $ax^2 + bx + c$  which is solved by completing the square – adding  $(b/2a)^2$  to each side. This is ‘generative’ reasoning, since the term  $(b/2a)^2$  is *not* in the term set we have defined. Equations that require these more exotic actions are beyond the scope of the current work (and were also not studied in all previous works (Poesia et al., 2021; Dabelow & Ueda, 2024)).

## 2.1. Inductive biases

We use a number of inductive biases to improve learning.

**TreeMLP.** We use a custom *TreeMLP* architecture to exploit the expression-tree structure of our equations (nodes are tokens/operators; edges are parent→child). Each node ID is embedded and linearly projected to a hidden size, then we perform  $K$  rounds of message passing with a simple sum aggregator over incoming edges (no degree normalization). At round  $t$ , a node’s representation is updated by applying a two-layer MLP with ReLU to the concatenation of its static embedding and the aggregated message, i.e.,  $h_i^{(t+1)} = \text{MLP}([\text{emb}_i \parallel \sum_{j \rightarrow i} h_j^{(t)}])$ . Padding-aware masks zero out invalid nodes and edges at every stage. After  $K$  iterations, we obtain a fixed-size graph feature by pooling (mean or max) over valid node states, which is fed to the policy/value heads. Compared to standard GCN/SAGE extractors, TreeMLP avoids degree normalization and learned edge transforms, yielding a lightweight, strictly local update that aligns well with symbolic expression trees and trains stably with PPO.

We give more details on TreeMLP in the Appendix. We also show this TreeMLP outperforms other popular graph based models, like GNNs and GraphSage (Figure 8).

**Curriculum learning.** We consider multi-equation environments, where the agent solves a single equation chosen randomly from a larger set during each episode. Equations are selected via *inverse sampling*, where the probability of choosing an equation is inversely proportional to the number of times it has been solved before. This curriculum ensures the agent spends more time trying to solve equations it hasn’t solved before.

**Success replay buffer.** A useful property of our environment is that solution traces are *exact*: once a sequence of

<sup>1</sup>We explored other values like  $|A| \in \{40, 70, 100\}$  and found no measurable change in performance.

<sup>2</sup>we judged number of nodes+edges correlates better with algebraic complexity than number of nodes alone.




Figure 2. Learning curves for the closed environments. Mean of three random seeds, shading from min to max.

operations solves an instance (e.g.,  $ax + b = 0$ ), replaying the same (op, term) actions deterministically reproduces the solution. We exploit this by enabling a light-weight *success replay* path in the env, where the solution traces  $(s_t, at)_t$  are stored to a memory buffer and periodically mixed into the on-policy rollouts.

**Relabel constants.** To improve parameter sharing across structurally identical equations that differ only by symbol names, we introduce a *relabel-constants* action. When invoked, it applies a consistent partial renaming (e.g.,  $\{d \mapsto a, e \mapsto b, \dots\}$ ) to the current main equation and the working state, mapping arbitrary coefficients onto a compact canonical alphabet  $\{a, b, c\}$ . The environment stores the induced substitution map on a stack (map\_constants\_history) and uses the mapped equation for all subsequent steps. Upon solving, we *unwind* the stack in reverse, substituting back the original symbols so the final answer is expressed in the user’s variables. This reduces input vocabulary entropy, prevents spurious overfitting to symbol IDs, and makes action terms like “subtract  $b$ ” reusable across many instances.

**Curiosity based rewards.** In previous work (Dabelow & Ueda, 2024), curiosity-based exploration bonuses (e.g. ICM, RND, NGU) were shown to enhance learning on a vanilla MLP policy. Stacking the same bonuses on top of our TreeMLP, however, *degrades* performance (Figure 10 in the appendix). We therefore do *not* use curiosity in our headline results.

Table 1. Greedy test accuracy for each agent across dataset, read from the end-of-training values of Figure 2. Values are approximate to within  $\pm 0.02$ . ConPoLe / ConPoLe-local numbers from (Poesia et al., 2021). Both PPO-TREE-RC and PPO-TREE-RC-BUF exceed the ConPoLe SOTA on poesia.

| Algorithm       | small       | large       | poesia      |
|-----------------|-------------|-------------|-------------|
| ConPoLe-local   | n/a         | n/a         | 0.765       |
| ConPoLe         | n/a         | n/a         | 0.925       |
| ppo             | 0.04        | 0.02        | 0.17        |
| ppo-tree        | 0.25        | 0.10        | 0.80        |
| ppo-tree-rc     | 0.85        | 0.10        | 0.85        |
| ppo-tree-buf    | 0.55        | 0.40        | 0.85        |
| ppo-tree-rc-buf | <b>0.85</b> | <b>0.50</b> | <b>0.93</b> |

### 3. Results: Closed equations

**Algorithms.** We use the standardized implementation of ppo from stable-baselines 3. We experimented with A2C and DQN but found worse performance. Hyperparameters, chosen via tuning, are reported in the Appendix. We trained for  $10^7$  steps and evaluated every  $5 \times 10^5$ . We used three random seeds.

**Datasets.** We tasked the agent with solving a random equation during each episode. We generate train/test sets by starting with  $x$  and applying actions recursively. Recall actions are (operation, term) pairs. The operations are as before, but we restrict terms to  $(a, b, c)$  (we excluded  $d, e$  since we wanted to favor long, rather than wide equations). Applying

a single action generated equations like  $ax$ ,  $\log(x)$ ,  $x+b$ , applying 2 actions equations like  $ax+b$ ,  $\log(x)+d$ ,  $(x+b)/c$  and so on. We threw out any equations that SymPy couldn't solve, and considered multi-root equations solved when just one root was found (e.g.  $\sin(x) = 0$  is solved if the agent found  $x = 0$  rather than  $x = 2n\pi$ ). We created a small dataset which included all equations of  $depth < 4$  (size 3874) and sub-sampled to a  $10^3/10^2$  train/test split, and a large dataset of all equations of  $depth < 5$  (size 15625) sub-sampled to a  $10^4/10^3$  train/test split. Finally, we also use the CommonCore datasets used in previous works.

**Results.** Figure 2 shows the learning curve for each algorithm on all three datasets. We plot three performance metrics: (i) coverage, the fraction of train equations solved at least once, (ii) test accuracy following the greedy policy, (iii) and test accuracy following beam search of width 5. The top row shows the CommonCore datasets (with coverage excluded since the train set is random and unbounded). The dashed black lines show the results for ConPoLe-local and ConPoLe from (Poesia et al., 2021), the SOTA for symbolic equation solving. On poesia, vanilla PPO-TREE reaches 0.74—comparable to ConPoLe-local (0.765)—and adding any single inductive bias pushes performance well above it: PPO-TREE-BUF attains 0.87, while the two best combinations exceed the full ConPoLe: PPO-TREE-RC reaches 0.96 and PPO-TREE-RC-BUF reaches **0.93**. Table 1 reports the exact numbers.

The bottom two rows show the results on our small and large datasets. For the small dataset, PPO-TREE-RC and PPO-TREE-RC-BUF achieve high test accuracy under greedy decoding ( $\geq 0.80$ ); the simpler variants PPO-TREE-BUF (0.47) and PPO-TREE (0.24) trail significantly. For the large dataset, only PPO-TREE-RC-BUF attains usable test accuracy (0.44 greedy,  $\text{test}_{\text{beam}} \approx 0.66$ ); vanilla PPO-TREE and PPO-TREE-RC collapse to near-zero on the large set, indicating that the success-replay buffer is the critical ingredient at scale.

The takeaway is that our inductive biases help in each case, and are comparable to the SOTA.

**Analysis.** How does the agent learn to solve the equation? We show two solution traces below. Notice the policy is not always optimal: the second example takes a 3-step detour (steps 5–7: expand, collect, truediv) where a single multiply-by- $-1$  would have sufficed, as in the first example's step 5.

1. Equation:  $cx + d + e(ax + b) = 0$

Step 1:  $cx + d + e(ax + b) = 0$  | expand, None  
Step 2:  $aex + be + cx + d = 0$  | collect, None  
Step 3:  $be + d + x*(ae + c) = 0$  | subtract,  $x(ae + c)$   
Step 4:  $-x(ae + c) = be + d$  | truediv,  $-(be + d)/(ae + c)$

Step 5:  $-x = (be + d)/(ae + c)$  | multiply,  $-1$   
Solved:  $x = -(be + d)/(ae + c)$

2. Equation:  $e + (ax + b)/(cx + d) = 0$

Step 1:  $e + (ax + b)/(cx + d) = 0$  | truediv,  $1/(cx + d)$   
Step 2:  $(e + (ax + b)/(cx + d))(cx + d) = 0$  | expand,  $(e + (ax + b)/(cx + d))(cx + d)$   
Step 3:  $ax + b + cex + de = 0$  | collect,  $x$   
Step 4:  $b + de + x*(a + ce) = 0$  | subtract,  $x(a + ce)$   
Step 5:  $-x(a + ce) = b + de$  | expand, None  
Step 6:  $-ax - cex = b + de$  | collect,  $x$   
Step 7:  $x*(-a - ce) = b + de$  | truediv,  $-(b + de)/(-a - ce)$   
Solved:  $x = (b + de)/(-a - ce)$

**Per-type accuracy.** The “small” dataset contains a heterogeneous mix of equation types. Bucketing the test set by structural class and re-evaluating the trained PPO-TREE-RC-BUF agent reveals a clear pattern (Table 2): linear and radical equations are nearly fully solved ( $\geq 0.88$ ), polynomial equations of degree 2–4 sit in the 0.75–0.82 range, while transcendental forms—particularly trigonometric and logarithmic—are the hardest. The trigonometric bucket (113 equations, 72% accuracy) accounts for most of the remaining failures. Note that the polynomial buckets contain only “depressed” forms (e.g.  $ax^2 + c$ ,  $x^4 + c$ ) that arise naturally from recursive action application; *true* open polynomials with a linear cross-term ( $ax^2 + bx + c$ ) are not in this dataset and are instead studied in §??.

Table 2. Greedy test accuracy of PPO-TREE-RC-BUF on the small dataset, broken down by equation class. “poly-degk” includes only the depressed forms reachable by the closed-equation action set (e.g.  $ax^2 + c$ ,  $x^4 + c$ ); equations with a linear cross-term are out of this dataset.

| Type                       | $n$ | solved | acc   |
|----------------------------|-----|--------|-------|
| radical (e.g. $\sqrt{x}$ ) | 47  | 43     | 0.915 |
| linear (poly-deg1)         | 110 | 97     | 0.882 |
| quadratic (poly-deg2)      | 34  | 28     | 0.824 |
| log                        | 77  | 58     | 0.753 |
| poly-deg4                  | 4   | 3      | 0.750 |
| trig                       | 113 | 81     | 0.717 |
| <b>overall</b>             | 387 | 312    | 0.806 |

**Learned representations.** As a qualitative probe of what the TreeMLP encodes, we extract the pooled graph embedding for each test equation and project it to two dimensions with t-SNE (Figure 3). Points are colored by structural type. Equations form well-separated clusters by family—linear, radical, trig, and log each occupy distinct regions—suggesting the representation captures structural identity rather than surface tokens.

**Failure modes.** Manual inspection of the failed cases (75 of 387) shows three recurring patterns: (i) trig equations that require an arccos/arcsin step the agent never selects, (ii)




Figure 3. t-SNE projection of TreeMLP pooled embeddings for 300 small-dataset equations, color-coded by structural type. Clusters correspond to equation family rather than to specific coefficient assignments.

log equations whose simplification creates a transcendental that breaks the closed-action set, and (iii) deep polynomials where the agent exceeds the per-episode action budget despite being on a productive path.

*Scaling to the large dataset.* The same per-type evaluation on the “large” dataset (1000 test equations of depth  $< 5$ ) reveals a uniform drop in every category: linear  $0.88 \rightarrow 0.60$ , quadratic  $0.82 \rightarrow 0.40$ , log  $0.75 \rightarrow 0.27$ , trig  $0.72 \rightarrow 0.36$ , radical  $0.92 \rightarrow 0.52$ . The drop is largest for deep polynomials (poly-deg4:  $0.75 \rightarrow 0.08$ ). Going from depth  $< 4$  to depth  $< 5$  approximately quadruples the test-set size and adds substantially deeper nesting, so the harder cases are genuinely harder, not just more numerous.

*Generalization to CommonCore.* On the CommonCore (Poesia) benchmark, the same PPO-TREE-RC agent reaches 0.955 greedy test accuracy. The breakdown is sharper: linear equations are solved at 0.98 and the small subset of rational equations at 1.00. The remaining failures are essentially edge-case templates that lack a free  $x$  in the polynomial decomposition (*degenerate* instances). Accounting for those, accuracy on the solvable subset is  $\geq 0.98$ .

## 4. Results: Open equations

Now we consider the harder class of *open equations*. We train a separate change-of-variables policy  $\pi_{\text{cov}}$  which, given an input equation, proposes a substitution  $x \mapsto \phi(x)$  that simplifies the instance into a canonical, easier-to-solve

form. We consider four basic templates

$$ax^2 + bx + c = 0, \quad x \rightarrow x - b/(2a) \quad (7)$$

$$ax^3 + bx^2 + cx + d = 0, \quad x \rightarrow x - b/(3a) \quad (8)$$

$$ax^4 + bx^3 + cx^2 + dx + e = 0, \quad x \rightarrow x - b/(4a) \quad (9)$$

$$ae^{kx} + be^{-kx} + c = 0, \quad x \rightarrow \frac{1}{k} \log x \quad (10)$$

The rhs of each shows the solving CoV. Rather than use a generative model (GAN/VAE) to propose  $\phi$ , we frame  $\pi_{\text{cov}}$  as an RL problem.

**MDP. States** are the encoded main equation together with the current step depth  $d$  and an action history (last  $H$  action indices). The encoded equation does not change during an episode; only  $d$  and the history evolve as the agent builds  $\phi$ . **Actions** are pairs  $(\text{op}, \tau)$  with  $\text{op} \in \{+, -, \times, \div\}$  and  $\tau$  drawn from a small term bank  $\{a, b, c, d, e, 2, 3, 4\}$ , plus a special STOP action; the first action fixes a *base op*, and subsequent actions build an inner substitution  $u$  compositionally. **Transitions** are deterministic: at termination (STOP or  $d = d_{\text{max}}$ ),  $\phi(x) := \text{base\_op}(x, u)$ , and the env substitutes  $x \mapsto \phi(x)$  in the main equation. **Reward** is the complexity-delta  $C(\text{main}) - C(\text{main after } \phi)$  with a small per-step penalty; large positive bonus when  $\phi$  reduces complexity (i.e., depresses the equation to a closed-equation-solvable form).

We train a CoV agent using PPO with TreeMLP. Figure 4 shows the results.




Figure 4. Performance of the CoV agent on open equations. The agent learns to propose correct substitutions to simplify the equations.

## Acknowledgements

### Impact Statement

This paper presents work whose goal is to advance the field of machine learning, specifically the application of reinforcement learning to symbolic mathematics. There are

---

many potential societal consequences of our work, none which we feel must be specifically highlighted here.

## References

- Dabelow, L. and Ueda, M. Symbolic equation solving via reinforcement learning. *arXiv preprint arXiv:2404.17499*, 2024.
- Degrave, J., Felber, F., Buchli, J., Neunert, M., Tracey, B., Carpanese, F., Ewalds, T., Hafner, R., Abdolmaleki, A., de Las Casas, D., et al. Magnetic control of tokamak plasmas through deep reinforcement learning. *Nature*, 602(7897):414–419, 2022.
- Langley, P. Crafting papers on machine learning. In *Proceedings of the 17th International Conference on Machine Learning (ICML 2000)*, pp. 1207–1216, Stanford, CA, 2000.
- Ng, A. Y., Coates, A., Diel, M., Ganapathi, V., Schulte, J., Tse, B., Berger, E., and Liang, E. Autonomous inverted helicopter flight via reinforcement learning. In *Experimental Robotics IX*, pp. 363–372. Springer, 2006.
- Poesia, G., Dong, W., and Goodman, N. Contrastive reinforcement learning of symbolic reasoning domains. In *Advances in Neural Information Processing Systems*, volume 34, pp. 17729–17741, 2021.
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and Klimov, O. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- Vinyals, O., Babuschkin, I., Czarnecki, W. M., Mathieu, M., Dudzik, A., Chung, J., Choi, D. H., Powell, R., Ewalds, T., Georgiev, P., et al. Grandmaster level in StarCraft II using multi-agent reinforcement learning. *Nature*, 575(7782):350–354, 2019.

## .1. Macroactions

Here we justify why the three macroactions (expand, None), (collect,  $x$ ), (mul,  $-1$ ) discussed in the main text are *required* to solve the equations we consider, not mere conveniences for the RL agent. Consider the action set *without* these macroactions:

$$O = \{\text{add, sub, mul, div}\} \cup \{\text{square, sqrt, exp, log, sin, cos, asin, acos}\}, \quad (11)$$

$$T = \text{SubExpr(lhs)} \cup \text{SubExpr(rhs)}, \quad (12)$$

$$A = (O \times T). \quad (13)$$

Now consider the equation  $dx + c(ax + b) = e$ . The agent must distribute  $c$  into the bracketed term  $ax + b$ , which the action set above cannot do: expand is required. Once expanded to  $dx + cax + cb = e$ , one needs to factor the  $x$  common to the first two terms—this is what (collect,  $x$ ) provides. Finally, consider  $-x = b$ : since  $-1$  is not in our term set, (mul,  $-1$ ) is needed to isolate  $x$ . (Alternatively,  $-1$  could be added to the term set, but we chose to keep the term set purely symbolic.)

## .2. Dataset generation: Rational Equation

To supplement the recursive datasets, we constructed rational equations of the form

$$\frac{ax + b}{cx + a} + b = 0,$$

where  $a, b, c$  are (non-zero) symbolic coefficients. We excluded degenerate cases such as denominators that vanish identically or cancel trivially with numerators.

Examples of retained equations include:

$$\begin{aligned} \frac{x + b}{cx + a} = 0 &\Rightarrow x = -b, \\ \frac{ax + b}{cx + a} + b = 0 &\Rightarrow x = \frac{-b - ab}{a + cb}, \\ \frac{ax + b}{cx + a} - c = 0 &\Rightarrow x = \frac{-b + ac}{a - c^2}. \end{aligned}$$

These rational forms were merged into both the small and large datasets, with their proportion limited to preserve diversity across other functional forms (logarithmic, trigonometric, exponential, etc.).

## .3. Illegal actions

The main issue was to avoid illegal division by zero. This is not as simple as precluding (div, 0) from the action set. We also had to prohibit *hidden* division by zero—for example, dividing by  $(x + a)$  or  $(x + b)$  when the equation has the form  $(x + a)(x + b) = 0$ , since one of the factors will vanish at the solution. To handle this, we check whether the current equation factors as  $P(x)Q(x) \cdots = 0$  with  $P, Q, \dots$  polynomial in  $x$ , and remove (div,  $P$ ), (div,  $Q$ ),  $\dots$  from the action set for that state.

## .4. Hyperparameters

We used Stable Baselines 3’s default hyperparameters for A2C and PPO (Table 3).

| Algorithm | Hyperparameters                                                                                                                                                                                                                                                                                                             |
|-----------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| PPO       | $\text{learning\_rate} = 0.0003, n\_steps = 2048, \text{batch\_size} = 64, n\_epochs = 10, \text{gamma} = 0.99, \text{gae\_lambda} = 0.95, \text{clip\_range} = 0.2, \text{ent\_coef} = 0.0, \text{vf\_coef} = 0.5, \text{max\_grad\_norm} = 0.5, \text{normalize\_advantage} = \text{True}, \text{device} = \text{'auto'}$ |
| PPO-tree  | $\text{learning\_rate} = 0.0007, n\_steps = 5, \text{gamma} = 0.99$                                                                                                                                                                                                                                                         |

Table 3. Default hyperparameters for A2C and PPO from Stable Baselines 3.

---

## .5. Solution Traces

Below are the solution traces for each equation in the fixed equation environment. Notice the solution trace is not always optimal. For instance,  $ax + b = 0$  could be solved with  $(sub, b)$ ,  $(div, a)$ , but instead the agent selects  $(sub, ax)$  and then  $(mul, -1)$ . "Truediv" is the sympy notation for divide.

1. Equation:  $a + x = 0$

```
Step 1: a + x = 0 | subtract, a
Solved: x = -a
```

2. Equation:  $ax = 0$

```
Step 1: a*x = 0 | div, a
Solved: x = 0
```

3. Equation:  $ax + b = 0$

```
Step 1: ax + b = 0 | subtract, ax
Step 2: -ax = b | multiply, -1
Step 3: ax = -b | div, a
Solved: x = -b/a
```

4. Equation:  $a/x + b = 0$

```
Step 1: a/x + b = 0 | subtract, b
Step 2: a/x = -b | truediv, 1/x
Step 3: -b*x = a | truediv, b
Step 4: -x = a/b | multiply, -1
Solved: x = -a/b
```

5. Equation:  $c(ax + b) + d = 0$

```
Step 1: c*(ax + b) + d = 0 | expand, None
Step 2: acx + bc + d = 0 | subtract, acx
Step 3: -acx = bc + d | multiply, -1
Step 4: acx = -bc - d | truediv, c
Step 5: ax = (-bc - d)/c | truediv, a
Solved: x = (-bc - d)/(ac)
```

6. Equation:  $c + d/(ax + b) = 0$

```
Step 1: c + d/(ax + b) = 0 | subtract, c
Step 2: d/(ax + b) = -c | multiply, (ax + b)
Step 3: d = -c (ax + b) | expand, None
Step 4: d = -c ax - c b | multiply, -1
Step 5: -d = c ax + c b | subtract, c b
Step 6: -d - c b = c ax | truediv, c
Step 7: (-d - c b)/c = ax | truediv, a
Step 8: ((-d - c b)/c)/a = x | truediv, 1/a
Solved: x = (-d - c b)/(c a)
```

7. Equation:  $cx + d + e(ax + b) = 0$

---

**Algorithm 1** TreeMLP message passing and pooling

---

**Require:** node IDs  $\mathbf{X}$ , edges (src, dst), masks  $m_{\text{node}}, m_{\text{edge}}$ , steps  $K$ 

- 1:  $\text{emb} \leftarrow \text{Proj}(\text{Embed}(\mathbf{X}))$   $\triangleright$  shape  $[B, N, H]$
  - 2:  $h \leftarrow \text{emb}$
  - 3: **for**  $k = 1$  **to**  $K$  **do**
  - 4:    $\mathcal{E} \leftarrow \{(i \rightarrow j) \mid m_{\text{edge}}(i \rightarrow j) = 1\}$
  - 5:    $\text{agg}[j] \leftarrow \sum_{(i \rightarrow j) \in \mathcal{E}} h[i]$   $\triangleright$  scatter-sum
  - 6:    $h \leftarrow \text{MLP}([\text{emb}, \text{agg}])$
  - 7:    $h \leftarrow h \odot m_{\text{node}}$   $\triangleright$  mask padded nodes
  - 8: **end for**
  - 9: **return**  $\text{pool}(h, m_{\text{node}})$   $\triangleright$  mean or max
- 

```
Step 1: cx + d + e(ax + b) = 0 | expand, None
Step 2: aex + be + cx + d = 0 | collect, x
Step 3: be + d + x*(ae + c) = 0 | subtract, x(ae + c)
Step 4: -x(ae + c) = be + d | truediv, ae + c
Step 5: -x = (be + d)/(ae + c) | multiply, -1
Solved: x = -(be + d)/(a*e + c)
```

8. Equation:  $e + (ax + b)/(cx + d) = 0$

```
Step 1: e + (ax + b)/(cx + d) = 0 | truediv, 1/(cx + d)
Step 2: (e + (ax + b)/(cx + d))(cx + d) = 0 | expand, None
Step 3: ax + b + cex + de = 0 | collect, x
Step 4: b + de + x*(a + ce) = 0 | subtract, x(a + ce)
Step 5: -x(a + ce) = b + de | expand, None
Step 6: -ax - cex = b + de | collect, x
Step 7: x*(-a - ce) = b + de | truediv, -a - ce
Solved: x = (b + de)/(-a - c*e)
```

## .6. TreeMLP Details

**Architecture.** TreeMLP embeds node IDs, projects to a hidden size, and performs  $K$  message-passing stages over the (directed) expression tree. At each stage it sums incoming neighbor states, concatenates the result with the fixed input embedding, and applies a two-layer MLP with ReLU. A masked global pool (mean or max) produces the graph embedding exposed to the policy/value heads.

*Listing 1.* TreeMLP forward pass (minimal).

```
def forward(self, obs: dict) -> torch.Tensor:
    node_ids = self._to_device(obs["node_features"], self.device)
    edge_index= self._to_device(obs["edge_index"], self.device)
    node_mask = self._to_device(obs["node_mask"], self.device)
    edge_mask = self._to_device(obs["edge_mask"], self.device)

    node_idx = self._id_to_idx(node_ids) if node_ids.dtype != torch.long else node_ids
    B, N = node_idx.shape
    _, _, E = edge_index.shape

    emb = self.embed(node_idx) # [B,N,embed_dim]
    emb = self.embed_proj(emb) # [B,N,H]
    h = emb.clone()

    src = edge_index[:, 0, :].long() # [B,E]
    dst = edge_index[:, 1, :].long() # [B,E]
    batch_idx = torch.arange(B, device=self.device).unsqueeze(1).expand(B, E)
```

---

```

for _ in range(self.K):
    src_flat = src.reshape(-1)
    dst_flat = dst.reshape(-1)
    b_flat = batch_idx.reshape(-1)
    m_flat = edge_mask.reshape(-1).bool()

    src_flat, dst_flat, b_flat = src_flat[m_flat], dst_flat[m_flat], b_flat[m_flat]
    src_idx = b_flat * N + src_flat
    dst_idx = b_flat * N + dst_flat

    h_flat = h.reshape(B * N, self.hidden_dim)
    messages = h_flat[src_idx] # [M, H]

    agg_flat = torch.zeros_like(h_flat) # [B*N, H]
    agg_flat.index_add(0, dst_idx, messages) # sum incoming
    agg = agg_flat.reshape(B, N, self.hidden_dim) # [B, N, H]

    h = self.node_mlp(torch.cat([emb, agg], dim=-1)) # [B, N, H]
    h = h * node_mask.unsqueeze(-1).float()

if self.pooling == "mean":
    denom = node_mask.sum(dim=1, keepdim=True).clamp(min=1).float()
    g = h.sum(dim=1) / denom
else:
    if self._minus_inf is None or self._dtype_cache != h.dtype:
        self._minus_inf = torch.finfo(h.dtype).min
        self._dtype_cache = h.dtype
    g = h.masked_fill(~node_mask.unsqueeze(-1), self._minus_inf).max(dim=1).values

return self.proj(g)

```

**Implementation notes.** We use fixed-size padding with boolean masks for nodes/edges, scatter-add for neighbor aggregation, and pool with mask awareness. All tensors are moved to CUDA when available; no MPS is used.

## .7. Representation visualizations

Figure 3 in the main text projects the TreeMLP pooled embedding to two dimensions using t-SNE. Here we show two additional views: a linear projection (PCA, more interpretable than t-SNE because it preserves global distance), and 3-dimensional versions of both projections (extra axis often reveals separation that the 2D projection collapses).

Figures 5, 6, 7 show the  $2 \times 2$  panel (PCA / t-SNE  $\times$  2D / 3D) on the small, large, and CommonCore (poesia) datasets respectively, color-coded by structural equation type. Three things are consistent across datasets and methods: (i) linear (poly-deg1) equations form a tight cluster; (ii) transcendental classes (trig, log) occupy distinct regions; (iii) PCA captures the same broad structure as t-SNE, indicating the family-level separation is not a t-SNE artifact.

## .8. Ablations

We summarize the ablation findings; the corresponding learning curves are in Figures 9 and 8.

- Replacing dense complexity-delta rewards with sparse outcome-only rewards (and disabling the inverse-frequency curriculum) substantially degrades performance. The main drawback of sparse rewards is *far* longer run times: we had to introduce a 1-second-per-step wall-clock timeout to keep training tractable.
- TreeMLP outperforms a vanilla GCN and GraphSAGE on the small dataset (Figure 8); the tree-aware aggregation appears necessary, not just helpful.
- Adding curiosity bonuses (ICM, RND, NGU) on top of TreeMLP *degrades* learning on `abel_level13` (Figure 10): at  $5 \times 10^5$  steps every curiosity variant trails the no-curiosity baseline on coverage and all three accuracy metrics, with RND failing to make any meaningful progress in 500k steps. This contrasts with earlier work (Dabelow & Ueda, 2024), where curiosity *improved* a vanilla MLP policy. A plausible explanation: the TreeMLP’s structural inductive

---

bias already supplies the exploration signal that curiosity provides for the MLP, making the two redundant or actively conflicting.




Figure 5. TreeMLP embeddings on the small dataset, projected by PCA (top) and t-SNE (bottom), in 2D (left) and 3D (right). 400 random equations, colored by structural type. The same cluster structure is visible across all four projections.




Figure 6. Same as Figure 5 but on the large dataset (depth  $< 5$ ). More types are present and cluster boundaries are softer, consistent with the lower per-type accuracies reported in the main text.




Figure 7. Same as Figure 5 but on CommonCore (poesia). The dataset is dominated by linear and rational equations, giving a simpler embedding geometry.




Figure 8. TreeMLP outperforms GCN and GraphSAGE on the small datasets. Mean of three random seeds with shading to min/max.




Figure 9. Ablation study on small datasets.




Figure 10. Curiosity bonuses (ICM, NGU, RND) degrade TreeMLP learning on `abel_level3`. All four panels (coverage, greedy, beam, success@10) at the same training budget ( $5 \times 10^5$  steps) show every curiosity variant tracking below the no-curiosity baseline. RND in particular fails to make meaningful progress. Baseline is mean  $\pm$  min/max across  $n=5$  seeds; curiosity variants are single seed.